

# CPSC 375 Project 1

Kenn Son, Hamid Suha, Vivian Truong

4/7/2022

```
setwd("~/Documents/CPSC-375")
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(class)
library(ggplot2)
```

1) Data preparation/wrangling to get all the data into one table that can be used for linear modeling

a)reading the data files using read\_csv()

```
covid <- read_csv("https://raw.githubusercontent.com/govex/COVID-19/master/data_tables/vaccine_data/global_data.csv")
```

```
## Rows: 831 Columns: 512
## -- Column specification -----
## Delimiter: ","
## chr   (5): iso2, iso3, Province_State, Country_Region, Combined_Key
## dbl  (505): UID, code3, Lat, Long_, Population, 2020-12-12, 2020-12-13, 2020-...
## lgl   (2): FIPS, Admin2
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
bed <- read_csv("hospitalbed.csv")
```

```
## Rows: 1770 Columns: 3
## -- Column specification -----
## Delimiter: ","
## chr (1): Country
## dbl (2): Year, Hospital beds (per 10 000 population)
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
demo <- read_csv("demographics.csv")
```

```
## Rows: 3885 Columns: 5
## -- Column specification -----
## Delimiter: ","
## chr (4): Country Name, Country Code, Series Name, Series Code
## dbl (1): YR2015
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

b) Removing unneeded rows (e.g., countries like Brazil and India report Province\_State-level data that is not needed as we are studying only country-level rates) and columns.

```
covid <- covid %>% filter(Population >= 0, is.na(Province_State))
covid <- covid %>% select(-Admin2, -FIPS, -Province_State, -UID, -iso2, -iso3, -code3, -Lat, -Long_, -C
covid
```

```
## # A tibble: 134 x 502
##   Country_Region Population '2020-12-12' '2020-12-13' '2020-12-14' '2020-12-15'
##   <chr>          <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 Afghanistan    38928341      NA           NA           NA           NA
## 2 Albania         2877800      NA           NA           NA           NA
## 3 Algeria        43851043       0            0            0            0
## 4 Andorra         77265        0            0            0            0
## 5 Angola         32866268      NA           NA           NA           NA
## 6 Antigua and B~  97928        NA           NA           NA           NA
## 7 Argentina      45195777       0            0            0            0
## 8 Australia      25459700      NA           NA           NA           NA
## 9 Austria        9006400       0            0            0            0
## 10 Azerbaijan    10139175      NA           NA           NA           NA
## # ... with 124 more rows, and 496 more variables: '2020-12-16' <dbl>,
## #   '2020-12-17' <dbl>, '2020-12-18' <dbl>, '2020-12-19' <dbl>,
## #   '2020-12-20' <dbl>, '2020-12-21' <dbl>, '2020-12-22' <dbl>,
## #   '2020-12-23' <dbl>, '2020-12-24' <dbl>, '2020-12-25' <dbl>,
## #   '2020-12-26' <dbl>, '2020-12-27' <dbl>, '2020-12-28' <dbl>,
## #   '2020-12-29' <dbl>, '2020-12-30' <dbl>, '2020-12-31' <dbl>,
## #   '2021-01-01' <dbl>, '2021-01-02' <dbl>, '2021-01-03' <dbl>, ...
```

```
demo <- demo %>% select(-`Series Name`, -`Country Code`)
```

c) tidying tables, as needed. For example, the vaccinations data is not tidy.

```
covid <- covid %>% pivot_longer(-c(Country_Region, Population), names_to = "Date", values_to = "shots")
covid <- covid %>% filter(shots > 0) %>% view()

#bed <- bed %>% arrange(Year) %>% pivot_wider( names_from = "Year",
#      values_from = "Hospital beds (per 10 000 population)")

demo <- demo %>% pivot_wider(names_from = "Series Code", values_from = YR2015)
```

d) Calculate the vaccination rate: vaccinations/population

```
covid <- covid %>% mutate(vacRate = shots/Population) %>% view()
covid
```

```
## # A tibble: 59,998 x 5
##   Country_Region Population Date      shots vacRate
##   <chr>          <dbl> <chr>    <dbl>    <dbl>
## 1 Afghanistan    38928341 2021-02-28 8200 0.000211
## 2 Afghanistan    38928341 2021-03-01 8200 0.000211
## 3 Afghanistan    38928341 2021-03-02 8200 0.000211
## 4 Afghanistan    38928341 2021-03-03 8200 0.000211
## 5 Afghanistan    38928341 2021-03-04 8200 0.000211
## 6 Afghanistan    38928341 2021-03-05 8200 0.000211
## 7 Afghanistan    38928341 2021-03-06 8200 0.000211
## 8 Afghanistan    38928341 2021-03-07 8200 0.000211
## 9 Afghanistan    38928341 2021-03-08 8200 0.000211
## 10 Afghanistan   38928341 2021-03-09 8200 0.000211
## # ... with 59,988 more rows
```

e) Since the most important factor affecting vaccination rate is the number of days since vaccination began (vaccination rate always increases), calculate a variable that is: number of days since first non-zero vaccination number. This variable will be important for modeling.

```
covid <- covid %>% group_by(Country_Region) %>% mutate(daysSinceStart = 1:n())
covid <- covid %>% select(-Date)
covid
```

```
## # A tibble: 59,998 x 5
## # Groups:   Country_Region [134]
##   Country_Region Population shots vacRate daysSinceStart
##   <chr>          <dbl> <dbl>    <dbl>          <int>
## 1 Afghanistan    38928341 8200 0.000211         1
## 2 Afghanistan    38928341 8200 0.000211         2
## 3 Afghanistan    38928341 8200 0.000211         3
## 4 Afghanistan    38928341 8200 0.000211         4
## 5 Afghanistan    38928341 8200 0.000211         5
## 6 Afghanistan    38928341 8200 0.000211         6
## 7 Afghanistan    38928341 8200 0.000211         7
## 8 Afghanistan    38928341 8200 0.000211         8
## 9 Afghanistan    38928341 8200 0.000211         9
## 10 Afghanistan   38928341 8200 0.000211        10
## # ... with 59,988 more rows
```

f) Discard data that is not needed. For example, only the number of hospital beds from the most recent year is necessary.

```
bed.new <- bed %>% group_by(Country) %>% summarise(Year = max(Year))
bed <- inner_join(bed.new, bed)
```

```
## Joining, by = c("Country", "Year")
```

```
bed <- bed %>% mutate(beds = `Hospital beds (per 10 000 population)` %>%
  select(-Year, -`Hospital beds (per 10 000 population)`)
```

```
bed
```

```
## # A tibble: 178 x 2
##   Country      beds
##   <chr>      <dbl>
## 1 Afghanistan    3.9
## 2 Albania       28.9
## 3 Algeria        19
## 4 Angola         8
## 5 Antigua and Barbuda 28.9
## 6 Argentina      49.9
## 7 Armenia       41.6
## 8 Australia      38.4
## 9 Austria       72.7
## 10 Azerbaijan   48.2
## # ... with 168 more rows
```

g) You can ignore sex-related differences in demographics in this project, so add the male/female population numbers together (already done in HW #5).

```
demo.total <- demo %>% mutate(SP.POP.80UP=SP.POP.80UP.FE+SP.POP.80UP.MA) %>%
  mutate(SP.POP.1564.IN=SP.POP.1564.MA.IN+SP.POP.1564.FE.IN) %>%
  mutate(SP.POP.0014.IN=SP.POP.0014.MA.IN+SP.POP.0014.FE.IN) %>%
  mutate(SP.DYN.AMRT=SP.DYN.AMRT.MA+SP.DYN.AMRT.FE) %>%
  mutate(SP.POP.TOTL.IN=SP.POP.TOTL.FE.IN+SP.POP.TOTL.MA.IN) %>%
  mutate(SP.POP.65UP.IN=SP.POP.65UP.FE.IN+SP.POP.65UP.MA.IN) %>%
  select(-contains(".FE")) %>% select(-contains(".MA"))
```

```
demo <- demo.total %>% group_by(`Country Name`) %>% summarise(SP.DYN.LE00.IN = sum(SP.DYN.LE00.IN, na.rm=T))
```

```
demo
```

```
## # A tibble: 259 x 10
##   'Country Name'      SP.DYN.LE00.IN SP.URB.TOTL SP.POP.TOTL SP.POP.80UP
##   <chr>              <dbl>      <dbl>      <dbl>      <dbl>
## 1 Afghanistan      63.4      8535606   34413603   85552
## 2 Albania          78.0     1654503   2880703    66965
## 3 Algeria          76.1     28146511  39728025   453741
## 4 American Samoa    0        48689    55812      0
## 5 Andorra           0        68919    78011      0
## 6 Angola           59.4     17691524  27884381   69363
```

```
## 7 Antigua and Barbuda          76.5      23392      93566      1571
## 8 Arab World                   71.2    229821020    396028278    2689793
## 9 Argentina                   76.1     39467043     43131966     1095211
## 10 Armenia                    74.5     1845585      2925553      77292
## # ... with 249 more rows, and 5 more variables: SP.POP.1564.IN <dbl>,
## #   SP.POP.0014.IN <dbl>, SP.DYN.AMRT <dbl>, SP.POP.TOTL.IN <dbl>,
## #   SP.POP.65UP.IN <dbl>
```

h) Merge all tables (Hint: Join using the country name)

#### *#Unifying Country Names*

```
demo <- demo %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Republic of Korea",
                                   "South Korea")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Egypt, Arab Rep.",
                                   "Egypt")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Gambia, The",
                                   "Gambia")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "St. Vincent and the Grenadines",
                                   "Saint Vincent and the Grenadines")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "St. Lucia",
                                   "Saint Lucia")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Lao PDR",
                                   "Laos")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Slovak Republic",
                                   "Slovakia")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Bahamas, The",
                                   "Bahamas")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Iran, Islamic Rep.",
                                   "Iran")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "Venezuela, RB",
                                   "Venezuela")) %>%
  mutate(`Country Name` = replace(`Country Name`, `Country Name` == "St. Kitts and Nevis",
                                   "Saint Kitts and Nevis"))

bed <- bed %>%
  mutate(Country = replace(Country, Country == "Venezuela (Bolivarian Republic of)",
                           "Venezuela")) %>%
  mutate(Country = replace(Country, Country == "Viet Nam",
                           "Vietnam")) %>%
  mutate(Country = replace(Country, Country == "United States",
                           "United States of America")) %>%
  mutate(Country = replace(Country, Country == "Iran (Islamic Republic of)",
                           "Iran")) %>%
  mutate(Country = replace(Country, Country == "Lao People's Democratic Republic",
                           "Laos")) %>%
  mutate(Country = replace(Country, Country == "Bolivia",
                           "Bolivia (Plurinational State of)"))

covid <- covid %>% mutate(Country_Region = replace(Country_Region,
                                                    Country_Region == "US", "United States of America"))

big_data <- covid %>% inner_join(bed, by=c(Country_Region = "Country")) %>%
  inner_join(demo, by=c(Country_Region = "Country Name"))
```

```
big_data <- big_data %>% rename(Country = Country_Region) %>%
  relocate(shots, .after = Country) %>% relocate(vacRate, .after = Country)

big_data

## # A tibble: 52,447 x 15
## # Groups:   Country [117]
##   Country      vacRate shots Population daysSinceStart  beds SP.DYN.LE00.IN
##   <chr>      <dbl> <dbl>      <dbl>          <int> <dbl>      <dbl>
## 1 Afghanistan 0.000211 8200    38928341           1  3.9        63.4
## 2 Afghanistan 0.000211 8200    38928341           2  3.9        63.4
## 3 Afghanistan 0.000211 8200    38928341           3  3.9        63.4
## 4 Afghanistan 0.000211 8200    38928341           4  3.9        63.4
## 5 Afghanistan 0.000211 8200    38928341           5  3.9        63.4
## 6 Afghanistan 0.000211 8200    38928341           6  3.9        63.4
## 7 Afghanistan 0.000211 8200    38928341           7  3.9        63.4
## 8 Afghanistan 0.000211 8200    38928341           8  3.9        63.4
## 9 Afghanistan 0.000211 8200    38928341           9  3.9        63.4
## 10 Afghanistan 0.000211 8200    38928341          10  3.9        63.4
## # ... with 52,437 more rows, and 8 more variables: SP.URB.TOTL <dbl>,
## #   SP.POP.TOTL <dbl>, SP.POP.80UP <dbl>, SP.POP.1564.IN <dbl>,
## #   SP.POP.0014.IN <dbl>, SP.DYN.AMRT <dbl>, SP.POP.TOTL.IN <dbl>,
## #   SP.POP.65UP.IN <dbl>
```

- 2) Linear modeling the Covid vaccination rate Make a list of all predictor variables that are available. The challenge is to identify which combination of these predictors will give the best predictive model. You should also try transforming some of the variables (e.g., transforming population counts to proportion of total population). Run linear regression with at least 5 different combinations of predictor variables. Note: each day becomes one data point, i.e., the vaccination rate is calculated for each day for each country. The number of vaccinations should not be used as an independent variable as this is essentially what you are predicting.

```
combo1 <- lm(vacRate~SP.POP.0014.IN+SP.POP.1564.IN+beds, data = big_data)
combo2 <- lm(vacRate~SP.DYN.AMRT+beds, data = big_data)
combo3 <- lm(vacRate~SP.POP.65UP.IN+beds+SP.POP.80UP, data = big_data)
combo4 <- lm(vacRate~SP.POP.80UP+beds, data = big_data)
combo5 <- lm(vacRate~Population+beds+daysSinceStart, data = big_data)

summary(combo1)

##
## Call:
## lm(formula = vacRate ~ SP.POP.0014.IN + SP.POP.1564.IN + beds,
##     data = big_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4162 -0.5805 -0.1556  0.5060  2.0604
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.535e-01  5.225e-03  125.07  <2e-16 ***
```

```
## SP.POP.0014.IN -6.597e-09 2.346e-10 -28.12 <2e-16 ***
## SP.POP.1564.IN 2.248e-09 7.994e-11 28.12 <2e-16 ***
## beds 3.688e-03 1.341e-04 27.49 <2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6636 on 52443 degrees of freedom
## Multiple R-squared: 0.04034, Adjusted R-squared: 0.04029
## F-statistic: 734.9 on 3 and 52443 DF, p-value: < 2.2e-16
```

```
summary(combo2)
```

```
##
## Call:
## lm(formula = vacRate ~ SP.DYN.AMRT + beds, data = big_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.15269 -0.52168 -0.06899  0.45815  1.89809
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.090e+00  7.489e-03 145.476 < 2e-16 ***
## SP.DYN.AMRT -1.408e-03  1.736e-05 -81.096 < 2e-16 ***
## beds         8.468e-04  1.309e-04   6.471 9.81e-11 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6304 on 52444 degrees of freedom
## Multiple R-squared: 0.1341, Adjusted R-squared: 0.1341
## F-statistic: 4063 on 2 and 52444 DF, p-value: < 2.2e-16
```

```
summary(combo3)
```

```
##
## Call:
## lm(formula = vacRate ~ SP.POP.65UP.IN + beds + SP.POP.80UP, data = big_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5546 -0.5897 -0.1543  0.5101  2.0911
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.294e-01  5.112e-03 123.12 <2e-16 ***
## SP.POP.65UP.IN -1.467e-08  1.084e-09 -13.53 <2e-16 ***
## beds          3.534e-03  1.487e-04  23.76 <2e-16 ***
## SP.POP.80UP    9.868e-08  6.395e-09  15.43 <2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6667 on 52443 degrees of freedom
## Multiple R-squared: 0.03139, Adjusted R-squared: 0.03134
## F-statistic: 566.6 on 3 and 52443 DF, p-value: < 2.2e-16
```

```
summary(combo4)
```

```
##
## Call:
## lm(formula = vacRate ~ SP.POP.80UP + beds, data = big_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3235 -0.5919 -0.1600  0.5082  2.0685
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 6.097e-01  4.909e-03  124.20  <2e-16 ***
## SP.POP.80UP 1.368e-08  1.190e-09   11.49  <2e-16 ***
## beds        4.484e-03  1.314e-04   34.13  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6679 on 52444 degrees of freedom
## Multiple R-squared:  0.02801, Adjusted R-squared:  0.02798
## F-statistic: 755.7 on 2 and 52444 DF, p-value: < 2.2e-16
```

```
summary(combo5)
```

```
##
## Call:
## lm(formula = vacRate ~ Population + beds + daysSinceStart, data = big_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.33445 -0.22383 -0.01228  0.24680  1.27957
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.644e-01  4.323e-03 -61.159  <2e-16 ***
## Population  -1.372e-11  9.554e-12  -1.436    0.151
## beds        3.905e-03  7.993e-05  48.853  <2e-16 ***
## daysSinceStart 4.004e-03  1.375e-05 291.148  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4134 on 52443 degrees of freedom
## Multiple R-squared:  0.6276, Adjusted R-squared:  0.6276
## F-statistic: 2.946e+04 on 3 and 52443 DF, p-value: < 2.2e-16
```

```
cf1 <- coef(combo1)
cf2 <- coef(combo2)
cf3 <- coef(combo3)
cf4 <- coef(combo4)
cf5 <- coef(combo5)
```

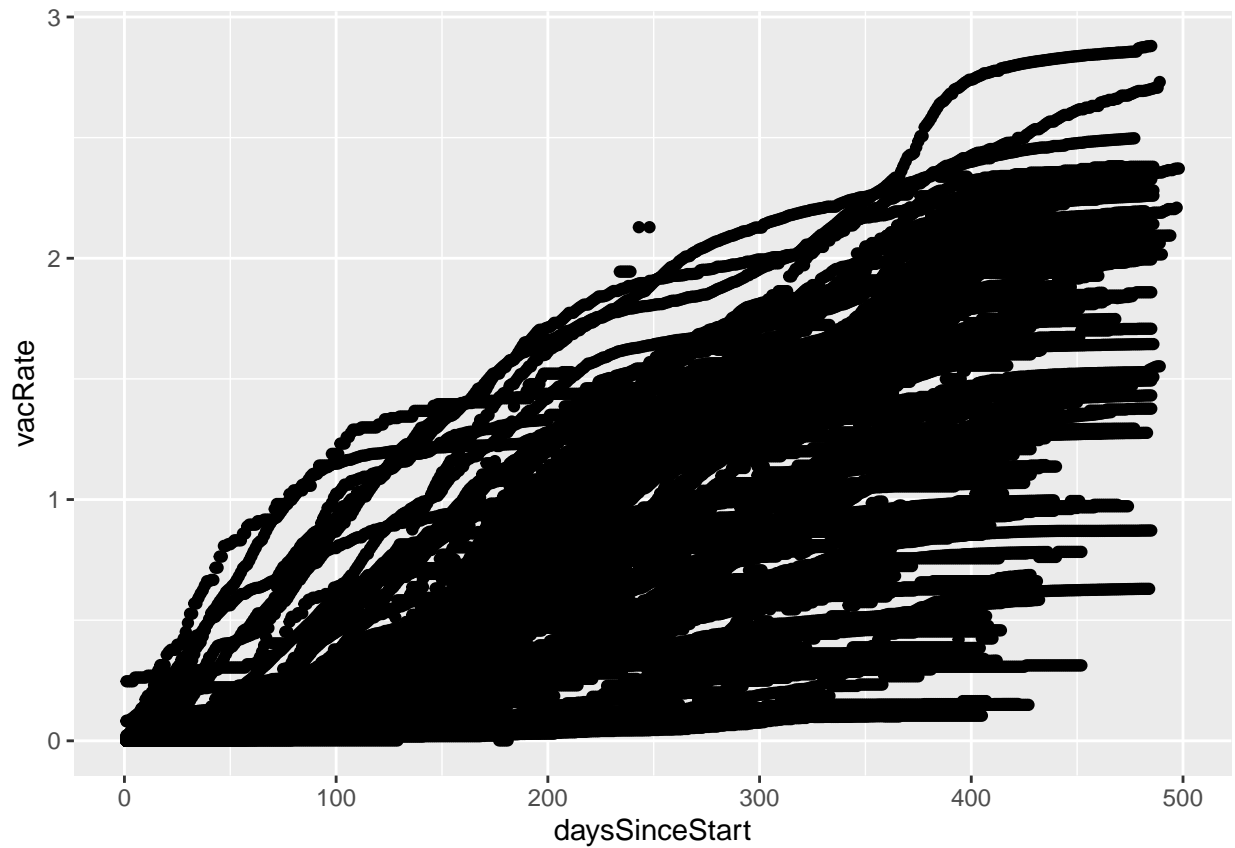
Write a short report describing your data wrangling steps and the different combinations of predictor variables you tried, and any variable transforms. [A PDF file]



The report should include the following plots:

i) a scatterplot of only the most recent vaccination rate for every country and the number of days since first vaccination

```
ggplot(big_data) + geom_point(mapping = aes(x = daysSinceStart, y = vacRate))
```



ii) a summary bar graph with the R2 values on the y-axis and a corresponding model name on the x-axis (include all the different models you tried).

```
R2 = c(0.04034,0.134,0.03136,0.028,0.6278)
big_data.R2 <- data.frame (Model = c("M1","M2","M3","M4","M5"),
                           Rsquared = R2)
ggplot(big_data.R2, aes(x = Model, y=Rsquared)) + geom_bar(stat = "identity", width=1)
```

